

SEC 1 Homework - 3

January 26, 2024

```
[15]: import pandas as pd
```

1 1.) Clean the Apple Data to get a quarterly series of EPS.

```
[16]: data = pd.read_csv('/Users/adrianonggowarsito/Desktop/AAPL_quarterly_financials.
↳csv', index_col=0, header=None).T
y = data.reindex(index=data.index[::-1])
y.set_index('name', inplace=True)
y.drop(y.tail(1).index, inplace=True)
y.index = pd.to_datetime(y.index)
Y = y['BasicEPS']
Y = Y.fillna(0.)
Y
```

```
[16]: name
1985-09-30      0.0
1985-12-31     0.004
1986-03-31     0.002
1986-06-30     0.002
1986-09-30      0.0
...
2022-09-30      1.29
2022-12-31      1.89
2023-03-31      1.53
2023-06-30      1.27
2023-09-30      1.47
Name: BasicEPS, Length: 153, dtype: object
```

2 1.) Come up with 6 search terms you think could nowcast earnings. (Different than the ones I used) Then, add in 3 terms that that you think will not Nowcast earnings. Pull in the gtrends data. Clean it to have a quarterly average.

```
[17]: from pytrends.request import TrendReq
```

```
[18]: import time
```

```
[19]: # Create pytrends object
pytrends = TrendReq(hl='en-US', tz=360)

# Set up the keywords and the timeframe
keywords = ['Phone', 'Sales', 'Macbook', 'Hack', 'OS', 'Tech', 'Food', 'Tokyo',
↳ 'Bali'] # Add your keywords here
start_date = '2004-01-01'
end_date = '2024-01-01'

# Create an empty DataFrame to store the results
df = pd.DataFrame()

# Iterate through keywords and fetch data
for keyword in keywords:
    time.sleep(5)
    pytrends.build_payload([keyword], cat=0, timeframe=f'{start_date}
↳ {end_date}', geo='', gprop='')
    interest_over_time_df = pytrends.interest_over_time()
    df[keyword] = interest_over_time_df[keyword]
```

```
[20]: dfq = df.resample('Q').mean()
dfq = dfq[dfq.index <= '2023-09-30']
dfq
```

```
[20]:
```

| | Phone | Sales | Macbook | Hack | OS | Tech \ |
|------------|-----------|-----------|-----------|-----------|-----------|-----------|
| date | | | | | | |
| 2004-03-31 | 80.333333 | 96.666667 | 0.000000 | 94.333333 | 67.333333 | 55.000000 |
| 2004-06-30 | 80.333333 | 91.000000 | 0.000000 | 83.666667 | 69.000000 | 49.000000 |
| 2004-09-30 | 85.666667 | 93.000000 | 0.000000 | 80.333333 | 61.333333 | 48.333333 |
| 2004-12-31 | 82.666667 | 81.666667 | 0.000000 | 80.000000 | 65.000000 | 47.666667 |
| 2005-03-31 | 80.333333 | 87.666667 | 0.000000 | 85.333333 | 62.000000 | 45.666667 |
| ... | ... | ... | ... | ... | ... | ... |
| 2022-09-30 | 91.000000 | 50.333333 | 83.000000 | 49.000000 | 80.333333 | 35.666667 |
| 2022-12-31 | 81.000000 | 48.333333 | 77.333333 | 42.000000 | 72.666667 | 36.333333 |
| 2023-03-31 | 76.333333 | 47.333333 | 76.000000 | 43.666667 | 65.000000 | 34.666667 |
| 2023-06-30 | 75.333333 | 49.000000 | 68.666667 | 43.333333 | 72.333333 | 31.666667 |
| 2023-09-30 | 79.333333 | 50.666667 | 78.000000 | 42.333333 | 65.000000 | 34.000000 |

| | Food | Tokyo | Bali |
|------------|-----------|-----------|-----------|
| date | | | |
| 2004-03-31 | 65.666667 | 14.666667 | 34.000000 |
| 2004-06-30 | 59.333333 | 14.666667 | 37.333333 |
| 2004-09-30 | 54.333333 | 14.000000 | 44.000000 |
| 2004-12-31 | 61.000000 | 14.333333 | 35.000000 |
| 2005-03-31 | 62.000000 | 13.666667 | 37.000000 |

```

...
2022-09-30  97.333333  14.666667  76.000000
2022-12-31  90.666667  14.333333  71.666667
2023-03-31  88.666667  17.000000  80.666667
2023-06-30  92.666667  17.000000  79.333333
2023-09-30  95.333333  17.000000  81.333333

```

[79 rows x 9 columns]

3 2.) Normalize all the X data

```
[21]: from sklearn.preprocessing import StandardScaler
```

```
[22]: scaler = StandardScaler()
```

```
[23]: X_scaled = scaler.fit_transform(dfq)
```

4 3.) Import data. Train, Test, Holdout (80%,15%,5%)

5 4.) Run a Lasso with lambda of .5. Plot a bar chart.

```
[24]: from sklearn.linear_model import Lasso
```

```
[25]: Y_04 = Y[Y.index >= '2004-03-31']
```

```
[26]: lasso = Lasso(alpha = .2) #lower lambda since its penalizing to much
```

```
[27]: lasso.fit(X_scaled, Y_04)
```

```
[27]: Lasso(alpha=0.2)
```

```
[28]: coefficients = lasso.coef_
```

```
[29]: import matplotlib.pyplot as plt
```

```

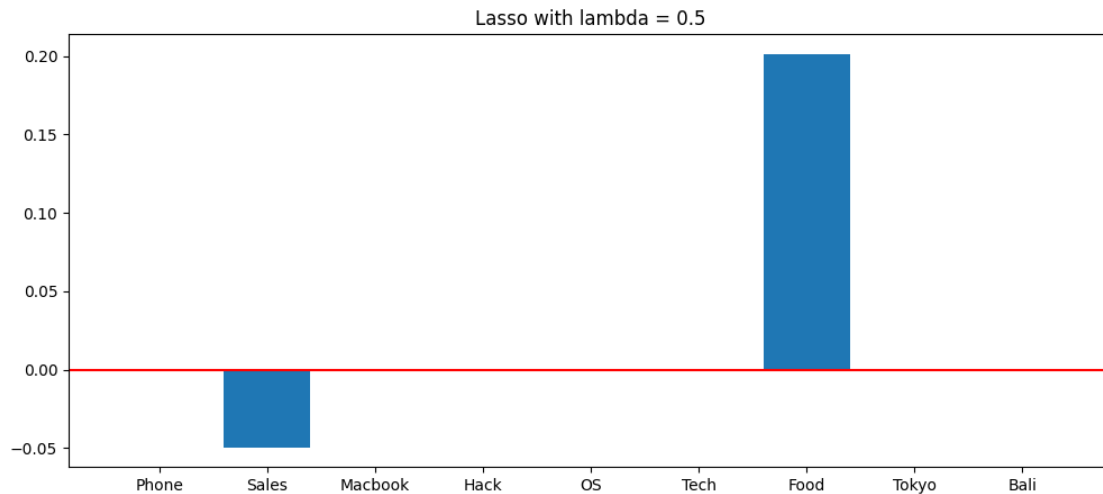
names = ['Phone', 'Sales', 'Macbook', 'Hack', 'OS', 'Tech', 'Food', 'Tokyo', '
↪ Bali']

```

```

plt.figure(figsize=(12,5))
plt.title('Lasso with lambda = 0.5')
plt.bar(range(len(coefficients)), coefficients, tick_label=names)
plt.axhline(0, color = 'red')
plt.show()

```



6 5.) Do these coefficient magnitudes make sense?

The coefficient does not really make sense. Much more relevant keywords like phone, macbook, os, and tech have 0 values, despite sales is affecting negatively, Food has pretty considerable outcome. Perhaps due to the amount of food bloggers using iPhones and/on social media

7 6.) Run a for loop looking at 10 different Lambdas and plot the coefficient magnitude for each.

```
[39]: import numpy as np

alphas = np.linspace(0, 0.5, 10)

# List to store coefficients for each alpha
coefficients_list = []

# Loop through different alpha values
for alpha in alphas:
    lasso = Lasso(alpha=alpha)
    lasso.fit(X_scaled, Y_04)
    coefficients = lasso.coef_
    coefficients_list.append(coefficients)

# Transpose the coefficients list for easier plotting
coefficients_array = np.array(coefficients_list).T

# Plotting
plt.figure(figsize=(12, 8))
```

```

for i in range(len(names)):
    plt.plot(alphas, coefficients_array[i], label=names[i])

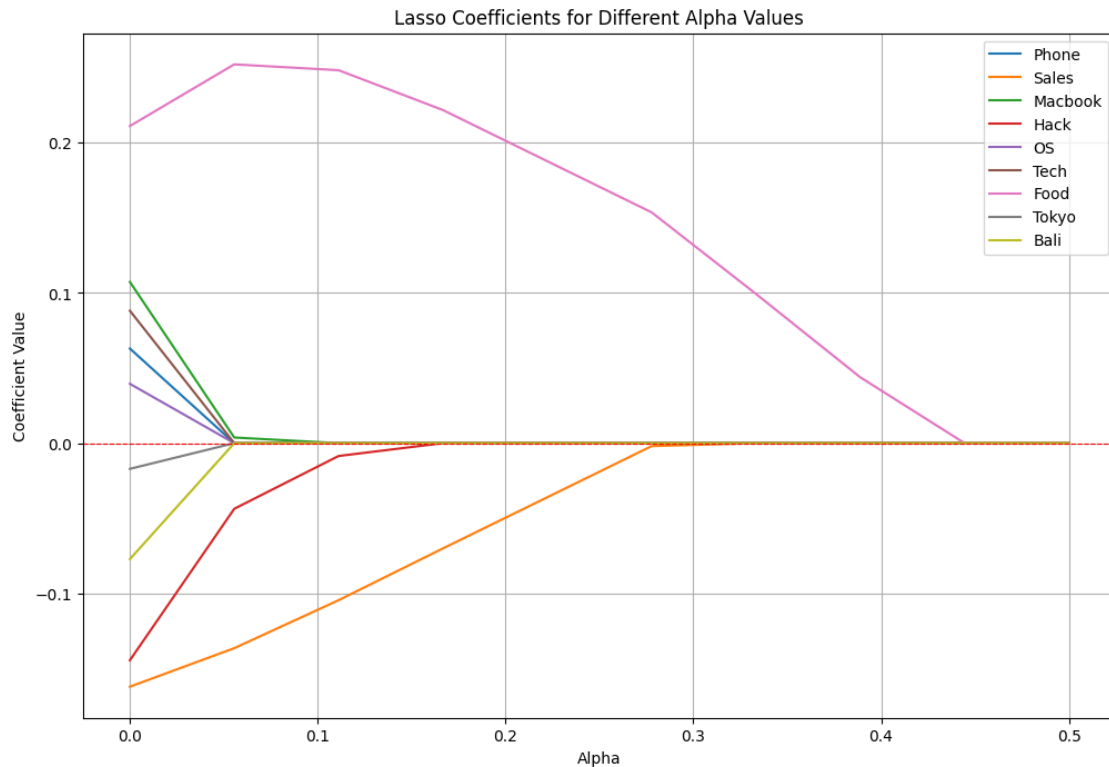
plt.xlabel('Alpha')
plt.ylabel('Coefficient Value')
plt.title('Lasso Coefficients for Different Alpha Values')
plt.axhline(0, color='red', linestyle='--', linewidth=0.8)
plt.legend()
plt.grid(True)
plt.show()

```

```

/Users/adrianonggowarsito/anaconda3/lib/python3.10/site-
packages/sklearn/base.py:1151: UserWarning: With alpha=0, this algorithm does
not converge well. You are advised to use the LinearRegression estimator
    return fit_method(estimator, *args, **kwargs)
/Users/adrianonggowarsito/anaconda3/lib/python3.10/site-
packages/sklearn/linear_model/_coordinate_descent.py:628: UserWarning:
Coordinate descent with no regularization may lead to unexpected results and is
discouraged.
    model = cd_fast.enet_coordinate_descent(
/Users/adrianonggowarsito/anaconda3/lib/python3.10/site-
packages/sklearn/linear_model/_coordinate_descent.py:628: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.945e+00, tolerance: 2.085e-03 Linear regression models with null weight
for the l1 regularization term are more efficiently fitted using one of the
solvers implemented in sklearn.linear_model.Ridge/RidgeCV instead.
    model = cd_fast.enet_coordinate_descent(

```

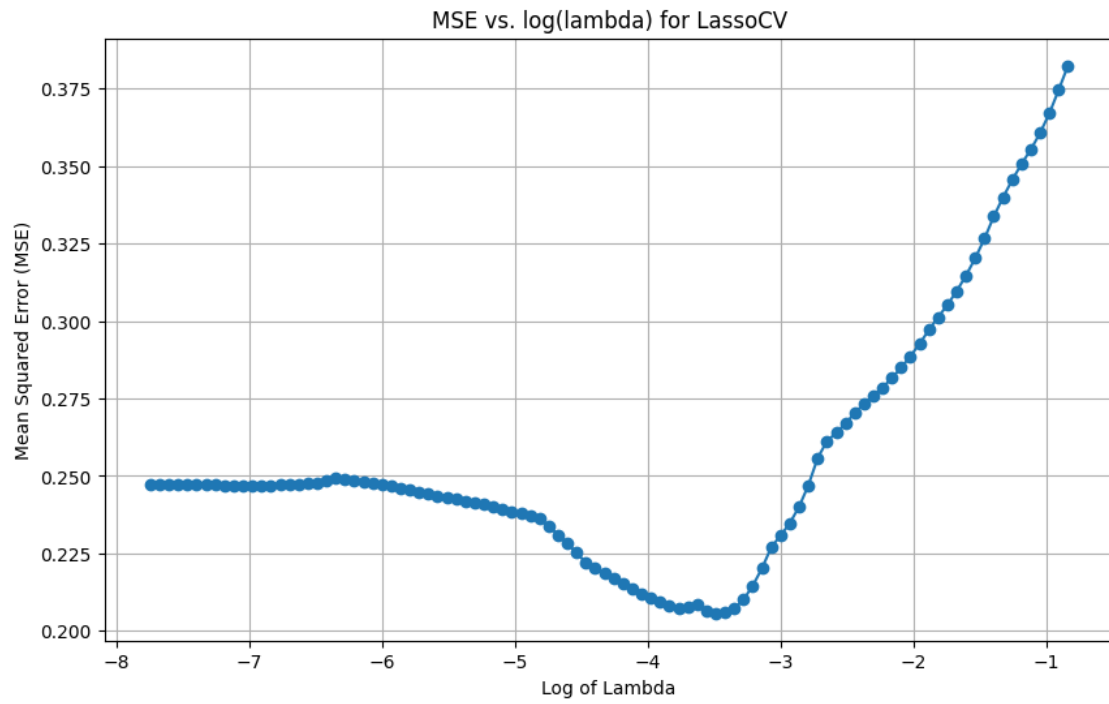


8 7.) Run a cross validation. What is your ideal lambda?

```
[38]: from sklearn.linear_model import LassoCV
modCV = LassoCV().fit(X_scaled, Y_04)

alphas = modCV.alphas_
mse_values = modCV.mse_path_.mean(axis=1)

# Plotting
plt.figure(figsize=(10, 6))
plt.plot(np.log(alphas), mse_values, marker='o')
plt.xlabel('Log of Lambda')
plt.ylabel('Mean Squared Error (MSE)')
plt.title('MSE vs. log(lambda) for LassoCV')
plt.grid(True)
plt.show()
```



8.0.1 Approximately between 3 and 4 (3.5)