

# Week 5 Coursework - Adrian

February 8, 2024

## 1 0.) Import the Credit Card Fraud Data From CCLE

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

```
[2]: df = pd.read_csv("/Users/adrianonggowarsito/Desktop/fraudTest.csv")
```

```
[3]: df.head()
```

```
[3]: Unnamed: 0 trans_date_trans_time      cc_num \
0          0  2020-06-21 12:14:25  2291163933867244
1          1  2020-06-21 12:14:33  3573030041201292
2          2  2020-06-21 12:14:53  3598215285024754
3          3  2020-06-21 12:15:15  3591919803438423
4          4  2020-06-21 12:15:17  3526826139003047

          merchant      category  amt  first \
0      fraud_Kirlin and Sons  personal_care  2.86  Jeff
1      fraud_Sporer-Keebler  personal_care  29.84  Joanne
2  fraud_Swaniawski, Nitzsche and Welch  health_fitness  41.28  Ashley
3      fraud_Haley Group      misc_pos  60.05  Brian
4      fraud_Johnston-Casper      travel   3.19  Nathan

          last gender      street  ...      lat      long \
0  Elliott      M      351 Darlene Green  ...  33.9659  -80.9355
1  Williams      F      3638 Marsh Union  ...  40.3207  -110.4360
2  Lopez      F      9333 Valentine Point  ...  40.6729  -73.5365
3  Williams      M  32941 Krystal Mill Apt. 552  ...  28.5697  -80.8191
4  Massey      M      5783 Evan Roads Apt. 465  ...  44.2529  -85.0170

          city_pop      job      dob \
0      333497  Mechanical engineer  1968-03-19
1          302  Sales professional, IT  1990-01-17
2      34496  Librarian, public  1970-10-21
3      54767  Set designer  1987-07-25
4          1126  Furniture designer  1955-07-06
```

	trans_num	unix_time	merch_lat	merch_long	\
0	2da90c7d74bd46a0caf3777415b3ebd3	1371816865	33.986391	-81.200714	
1	324cc204407e99f51b0d6ca0055005e7	1371816873	39.450498	-109.960431	
2	c81755dbbbea9d5c77f094348a7579be	1371816893	40.495810	-74.196111	
3	2159175b9efe66dc301f149d3d5abf8c	1371816915	28.812398	-80.883061	
4	57ff021bd3f328f8738bb535c302a31b	1371816917	44.959148	-85.884734	

	is_fraud
0	0
1	0
2	0
3	0
4	0

[5 rows x 23 columns]

```
[4]: df_select = df.copy()
```

```
[5]: df_select = df_select[
      ["trans_date_trans_time", "category", "amt", "city_pop", "is_fraud"]
    ]
```

```
[6]: df_select["trans_date_trans_time"] = pd.to_datetime(
      df_select.trans_date_trans_time
    )
```

```
[7]: df_select["time_var"] = [
      i.second for i in df_select["trans_date_trans_time"]
    ]
```

```
[8]: X = pd.get_dummies(df_select, columns=["category"])
      X = X.drop(["trans_date_trans_time", "is_fraud"], axis = 1)
```

```
[9]: y = df["is_fraud"]
```

## 2 1.) Use scikit learn preprocessing to split the data into 70/30 in out of sample

```
[10]: from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler
```

```
[11]: # Splitting the data between training and testing.
      X_train, X_test, y_train, y_test = train_test_split(
          X, y,
          test_size=0.3,
```

```

    random_state=999
)

```

```

[12]: # Splitting the testing data between testing and holdout.
X_test, X_holdout, y_test, y_holdout = train_test_split(
    X_test, y_test,
    test_size=0.5
)

```

### 3 2.) Make three sets of training data (Oversample, Undersample and SMOTE)

```

[13]: from imblearn.over_sampling import RandomOverSampler
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression

```

```

[14]: # Create the oversampled training data
over_sampler = RandomOverSampler(random_state=999)
X_over, y_over = over_sampler.fit_resample(X_train, y_train)

# Create the undersampled training data
under_sampler = RandomUnderSampler(random_state=999)
X_under, y_under = under_sampler.fit_resample(X_train, y_train)

# Apply SMOTE to your training data
smote = SMOTE(random_state=999)
smote_X, smote_y = smote.fit_resample(X_train, y_train)

```

### 4 3.) Train three logistic regression models

```

[15]: # Define the pipeline for the logistic regression model with oversampling
log_over = Pipeline([
    ('scaler', StandardScaler()),
    ('model', LogisticRegression(random_state=999))
])
# Fit the model on the oversampled dataset
log_over.fit(X_over, y_over)

```

```

[15]: Pipeline(steps=[('scaler', StandardScaler()),
    ('model', LogisticRegression(random_state=999))])

```

```

[16]: # Define the pipeline for the logistic regression model with undersampling
log_under = Pipeline([

```

```

        ('scaler', StandardScaler()),
        ('model', LogisticRegression(random_state=999))
    ])
    # Fit the model on the undersampled dataset
    log_under.fit(X_under, y_under)

```

```

[16]: Pipeline(steps=[('scaler', StandardScaler()),
                      ('model', LogisticRegression(random_state=999))])

```

```

[17]: # Define the pipeline for the logistic regression model with SMOTE
log_smote = Pipeline([
    ('scaler', StandardScaler()),
    ('balancer', SMOTE(random_state=999)),
    ('model', LogisticRegression(random_state=999))
])
# Fit the model on the original training dataset since SMOTE is part of the
↳ pipeline
log_smote.fit(X_train, y_train)

```

```

[17]: Pipeline(steps=[('scaler', StandardScaler()),
                      ('balancer', SMOTE(random_state=999)),
                      ('model', LogisticRegression(random_state=999))])

```

## 5 4.) Test the three models

```

[18]: def print_scores(x, y, over, under, smote, title = 'Out of Sample'):
    # Calculating the score
    over = over.score(x, y)
    under = under.score(x, y)
    smote = smote.score(x, y)

    temp = f"""
    Test Scores for {title}
    Accuracy
    -----
    - Over Sample: {over:.4f}
    - Under Sample: {under:.4f}
    - SMOTE: {smote:.4f}
    """
    print(temp)

```

Testing the three models in-sample.

```

[19]: print_scores(
    X_train, y_train,
    log_over, log_under, log_smote,
    title = 'In-Sample'
)

```

```
)
```

```
Test Scores for In-Sample
Accuracy
-----
- Over Sample: 0.9126
- Under Sample: 0.8948
- SMOTE: 0.9067
```

Testing the three models out-sample.

```
[20]: print_scores(
      X_test, y_test,
      log_over, log_under, log_smote,
      title = 'Out-Sample'
    )
```

```
Test Scores for Out-Sample
Accuracy
-----
- Over Sample: 0.9116
- Under Sample: 0.8949
- SMOTE: 0.9062
```

```
[21]: # We see SMOTE performing with higher accuracy but is ACCURACY really the best_
      ↪measure?
```

Accuracy is not the best measure for this dataset. We are more concerned about to detect fraud, so we want to maximize the number of true positives and minimize the number of false negatives. In other words, we want to maximize the sensitivity (recall) and minimize the false negative rate.

## 6 5.) Which performed best in Out of Sample metrics?

```
[22]: print_scores(
      X_holdout, y_holdout,
      log_over, log_under, log_smote,
      title = 'HoldOut-Sample'
    )
```

```
Test Scores for HoldOut-Sample
Accuracy
-----
- Over Sample: 0.9138
- Under Sample: 0.8959
```

- SMOTE: 0.9079

Analyzing the out of sample and holdout accuracy, we see that oversample is better, followed by SMOTE.

```
[23]: # Sensitivity here in credit fraud is more important as seen from last class
```

```
[24]: from sklearn.metrics import confusion_matrix
```

```
[25]: y_true = y_test
```

```
[26]: # Assuming y_test is already defined  
  
# Over-sampling model predictions and sensitivity  
y_pred_over = log_over.predict(X_test)  
cm_over = confusion_matrix(y_test, y_pred_over)  
print("Over Sample Sensitivity : ", cm_over[1,1] / (cm_over[1,0] +  
    ↪cm_over[1,1]))  
  
# Under-sampling model predictions and sensitivity  
y_pred_under = log_under.predict(X_test)  
cm_under = confusion_matrix(y_test, y_pred_under)  
print("Under Sample Sensitivity : ", cm_under[1,1] / (cm_under[1,0] +  
    ↪cm_under[1,1]))  
  
# SMOTE model predictions and sensitivity  
y_pred_smote = log_smote.predict(X_test)  
cm_smote = confusion_matrix(y_test, y_pred_smote)  
print("SMOTE Sample Sensitivity : ", cm_smote[1,1] / (cm_smote[1,0] +  
    ↪cm_smote[1,1]))
```

Over Sample Sensitivity : 0.70625

Under Sample Sensitivity : 0.70625

SMOTE Sample Sensitivity : 0.709375


## 7 6.) Pick two features and plot the two classes before and after SMOTE.

```
[27]: raw_temp = pd.concat([X_train, y_train], axis =1)
```

```
[28]: #plt.scatter(raw_temp[raw_temp["is_fraud"] == 0]["amt"],  
    ↪raw_temp[raw_temp["is_fraud"] == 0]["city_pop"])  
  
plt.scatter(raw_temp[raw_temp["is_fraud"] == 1]["amt"],  
    ↪raw_temp[raw_temp["is_fraud"] == 1]["city_pop"])  
plt.legend(["Fraud", "Not Fraud"])
```

```
plt.xlabel("Amount")
plt.ylabel("Population")

plt.show()
```



```
[29]: raw_temp = pd.concat([smote_X, smote_y], axis=1)
```


```
[30]: #plt.scatter(raw_temp[raw_temp["is_fraud"] == 0]["amt"],
#               ↪raw_temp[raw_temp["is_fraud"] == 0]["city_pop"])

plt.scatter(raw_temp[raw_temp["is_fraud"] == 1]["amt"],
            ↪raw_temp[raw_temp["is_fraud"] == 1]["city_pop"])
plt.legend([ "Not Fraud", "Fraud"])
plt.xlabel("Amount")
plt.ylabel("Population")

plt.show()
```

/Users/adrianonggowarsito/anaconda3/lib/python3.10/site-packages/IPython/core/pylabtools.py:152: UserWarning: Creating legend with

loc="best" can be slow with large amounts of data.  
fig.canvas.print\_figure(bytes\_io, \*\*kw)





- 8 7.) We want to compare oversampling, Undersampling and SMOTE across our 3 models (Logistic Regression, Logistic Regression Lasso and Decision Trees).
- 9 Make a dataframe that has a dual index and 9 Rows.
- 10 Calculate: Sensitivity, Specificity, Precision, Recall and F1 score. for out of sample data.
- 11 Notice any patterns across performance for this model. Does one totally out perform the others IE. over/under/smote or does a model perform better DT, Lasso, LR?
- 12 Choose what you think is the best model and why. test on Holdout

```
[31]: from sklearn.tree import DecisionTreeClassifier
      from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score
      import pandas as pd
```

```
[32]: model_configs = {
      'log': LogisticRegression(),
      'lasso': LogisticRegression(
          penalty = 'l1', C = 0.5, solver = 'liblinear'
      ),
      'tree': DecisionTreeClassifier()
    }

    balancing_configs = {
      'over': RandomOverSampler(),
      'under': RandomUnderSampler(),
      'smote': SMOTE()
    }
```

```
[33]: trained_models = {}
      scores_for_df = {}
```

```
[34]: for i,j in balancing_configs.items():
      for k,l in model_configs.items():
          pipe = Pipeline(
              steps=[
                  ('scaler', StandardScaler()),
                  ('balancer', j),
                  ('model', l)
              ]
          )
```

```

    ]
)
pipe.fit(X_train, y_train)
trained_models[(i,k)] = pipe
# Compute precision, recall, f1 score and store them in a dictionary
y_pred = pipe.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
sensitivity = cm[1,1] / ( cm[1,0] + cm[1,1])
specificity = cm[0,0] / ( cm[0,0] + cm[0,1])
accuracy = pipe.score(X_test, y_test)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
scores_for_df[(i,k)] = {
    'precision': precision,
    'recall': recall,
    'sensitivity': sensitivity,
    'specificity': specificity,
    'accuracy': accuracy,
    'f1': f1,
}

```

```
[35]: pd.DataFrame(scores_for_df).T
```

```
[35]:
```

		precision	recall	sensitivity	specificity	accuracy	f1
over	log	0.029562	0.706250	0.706250	0.910655	0.909871	0.056748
	lasso	0.029817	0.709375	0.709375	0.911053	0.910279	0.057229
	tree	0.584615	0.475000	0.475000	0.998699	0.996689	0.524138
under	log	0.030534	0.712500	0.712500	0.912823	0.912054	0.058559
	lasso	0.034721	0.706250	0.706250	0.924336	0.923499	0.066188
	tree	0.065563	0.928125	0.928125	0.949023	0.948943	0.122474
smote	log	0.029006	0.709375	0.709375	0.908488	0.907723	0.055733
	lasso	0.029573	0.709375	0.709375	0.910294	0.909523	0.056778
	tree	0.259080	0.668750	0.668750	0.992630	0.991387	0.373473

Looking at the provided scores, if we prioritize F1 score (a balanced metric for precision and recall), the Decision Tree classifier with undersampling (“under\_tree”) seems to perform the best. It has the highest F1 score, indicating a good balance between precision and recall. It also has the highest recall, which is crucial for fraud detection as it’s more important to catch as many frauds as possible.

However, the trade-off is that it has a lower specificity compared to the other models, which means it is more likely to incorrectly label non-fraudulent transactions as fraudulent. Depending on the cost and impact of false positives, this might be an acceptable trade-off.

On the other hand, if we want to maintain higher specificity (reducing false positives), SMOTE has a slightly lower recall but better specificity than undersampling.